

INFO 210: Database Management Systems

Homework 1 solutions

- This assignment covers the following topics
 - Set operation
 - The relational model

Problem 1: Sets

Let us denote by **P** the set of all **people**, by **S** the set of people who are **single**, by **C** the set of people who have **children**, by **W** the set of **women** and by **M** the set of **men**. For each question below, write down an expression that represents the set being described. Write exactly **one expression** for each question.

P --- set of all people,

S --- set of people who are single,

C --- set of people who have children,

W --- set of women, and

M --- set of men

The usual **set operations**:

union \cup

intersection \cap

set difference \setminus

Problem 1: Sets

P --- set of all people,
 S --- set of people who are **single**,
 C --- set of people who **have children**,
 W --- set of **women**, and
 M --- set of **men**

The usual **set operations**:

union \cup

intersection \cap

set difference \setminus

Write exactly one expression for each question.

- a) Single men who do not have children.
 – $(S \cap M) - C$
- b) Single women who do not have children.
 – $(S \cap W) - C$
- c) The number of women who have children.
 – $|W \cap C - S|$ $| (W - S) \cap C |$
- d) Women who are either married with children, or single and do not have children.
 – $((W - S) \cap C) \cup ((W \cap S) - C)$ $((W \cap C) - S) \cup ((S \cap W) - C)$
- e) All possible married heterosexual couples, i.e., all possible pairings of married men and married women.
 – $(M - S) \times (W - S)$

Problem 2 : Keys

- Consider an entity set *Person*, with attributes social security number (ssn), name, nickname, address, and date of birth (dob). Assume that the following conditions hold:
 - (1) no two persons have the same ssn;
 - (2) no two persons have the same combination of name, address, and dob.
 - Further, assume that all persons have an ssn, a name and a dob, but that some persons don't have a nickname nor an address.
- **(a)** List **all** candidate keys and **all** superkeys for this entity set.
- **(b)** Write a create table statement that defines a relation appropriate for this entity set.

Problem 2 : Keys

- Consider an entity set *Person*, with attributes social security number (**ssn**), **name**, **nickname**, **address**, and date of birth (**dob**).
- (a) List **all** candidate keys and **all** superkeys for this entity set.

There are 2 candidate keys: (**ssn**) and (**name, address, dob**).

First, we list all sets of attributes that include the first candidate key, (**ssn**), as a subset.

(ssn, name), (ssn, nickname), (ssn, address), (ssn, dob),

(ssn, name, nickname), (ssn, name, address), (ssn, name, dob),
(ssn, nickname, address), (ssn, nickname, dob), (ssn, address, dob),

(ssn, name, nickname, address), (ssn, name, nickname, dob),

(**ssn, name, address, dob**), (ssn, nickname, address, dob),

(**ssn, name, nickname, address, dob**)

There are exactly $4 + 6 + 4 + 1 = 15$ such superkeys.

Problem 2 : Keys

- Consider an entity set *Person*, with attributes social security number (**ssn**), **name**, **nickname**, **address**, and date of birth (**dob**).
- (a) List **all** candidate keys and **all** superkeys for this entity set.

There are 2 candidate keys: (ssn) and (name, address, dob).

Next, we list all sets of attributes that include the second candidate key, (name, address, dob), as a subset. There are exactly $2+1=3$ such superkeys.

(**ssn, name, address, dob**), (name, nickname, address, dob),

(**ssn, name, nickname, address, dob**)

Note that 2 superkeys (in bold) include both candidate keys as subsets.

Thus, there are a total of $15 + 3 - 2 = 16$ superkeys for this relation.

Problem 2 : Keys

- Consider an entity set *Person*, with attributes social security number (ssn), name, nickname, address, and date of birth (dob). Assume that the following conditions hold:
 - (1) no two persons have the same ssn;
 - (2) no two persons have the same combination of name, address, and dob.
 - Further, assume that **all** persons **have** an ssn, a name and a dob, but that **some** persons **don't have** a nickname nor an address.
- (b)** Write a create table statement that defines a relation appropriate for this entity set.
(Lengths of the varchar in this solution could be any the proper values)

```
create table Person (  
    ssn char (11) primary key,  
    name varchar(64),  
    nickname varchar(32),  
    address varchar(128),  
    dob date,  
    unique (name, address, dob)  
);
```


Problem 2 : Keys

- Consider an entity set *Person*, with attributes social security number (ssn), name, nickname, address, and date of birth (dob). Assume that the following conditions hold:
 - (1) no two persons have the same ssn;
 - (2) no two persons have the same combination of name, address, and dob.
 - Further, assume that **all** persons **have** an ssn, a name and a dob, but that **some** persons **don't have** a nickname nor an address.
- (b)** Write a create table statement that defines a relation appropriate for this entity set.

```
create table Person (  
  ssn char (11) unique,  
  name varchar(64),  
  nickname varchar(32),  
  address varchar(128),  
  dob date,  
  primary key (name, address, dob)  
);
```

~~try to designate ssn as unique, and (name, address, dob) as a primary key.
But primary key requires “unique” and “not null” constraint **By Default!**~~

Problem 3: Schemas and instances

Consider an instance of the relation *Foo*. Below, we ask you to write three create table statements. Each create table statement must define a primary key.

Foo (A, B, C, D)

A	B	C	D
1	Ann	23	3
2	Bob	23	4
3	Joe	20	3
4	Bob	20	4

(a) Write two different create table statements for which the instance of *Foo* is legal. Note that taking the first statement and simply reordering columns does not give a different create table statement.

(b) Write a create table statement that would make the instance of *Foo* above illegal.

Problem 3: Schemas and instances

A	B	C	D
1	Ann	23	3
2	Bob	23	4
3	Joe	20	3
4	Bob	20	4

(a) Write **two different** create table statements for which the instance of *Foo* is **legal**. Note that taking the first statement and simply reordering columns does not give a different create table statement.

```
create table Foo (  
  A      number,  
  B      char(3),  
  C      number,  
  D      number,  
  primary key (C,D)  
);
```

```
create table Foo (  
  B      char(3),  
  A      number,  
  D      number,  
  C      number,  
  primary key (C,D)  
);
```

```
create table Foo (  
  A      number primary key,  
  B      char(3),  
  C      number,  
  D      number  
);
```

Several other solutions are possible for question (a) and (b).

Problem 3: Schemas and instances

A	B	C	D
1	Ann	23	3
2	Bob	23	4
3	Joe	20	3
4	Bob	20	4

(b) Write a create table statement that would make the instance of *Foo* above **illegal**.

```
create table Foo (  
    A    number,  
    B    char(3) primary key,  
    C    number,  
    D    number  
);
```

Several other solutions are possible for question (a) and (b).

Problem 4: Foreign keys

- Consider relation schemas below, with primary keys underlined.
 - City (name, state, population, elevation)
 - State (name, region)
 - Mayor (name, city, state, party)
 - Governor (name, state, party)
- Suppose that the following business rules hold. Each mayor governs exactly one city. Each governor governs exactly one state.
- **(a)** Write create table statements that encode these relation schemas and business rules with the right foreign key constraints.
- **(b)** In what order would you drop these tables? Give **all** valid sequences.

Problem 4: Foreign keys

- Consider relation schemas below, with primary keys underlined.
 - City (name, state, population, elevation)
 - State (name, region)
 - Mayor (name, city, state, party)
 - Governor (name, state, party)
- Each **mayor governs** exactly one **city**. Each **governor governs** exactly one **state**.
- (a)** Write create table statements that encode these relation schemas and business rules with the right foreign key constraints.

```
create table City (  
    name      varchar(64),  
    state     varchar(32),  
    population number,  
    elevation  number,  
    primary key (name, state)  
);
```


```
create table State (  
    name varchar(32) primary key,  
    region varchar(32)  
);
```

```
create table Mayor (  
    name varchar(128),  
    city  varchar(64),  
    state varchar(32),  
    party varchar(32),  
    primary key (name, city, state),  
    foreign key (city, state) references City(name, state)  
);
```

```
create table Governor (  
    name varchar(128),  
    state varchar(32),  
    party varchar(32),  
    primary key (name, state),  
    foreign key (state) references State(name)  
);
```

Problem 4: Foreign keys

- Consider relation schemas below, with primary keys underlined.
 - City (name, state, population, elevation)
 - State (name, region)
 - Mayor (name, city, state, party)
 - Governor (name, state, party)
- Each mayor governs exactly one city. Each governor governs exactly one state.
- **(b)** In what order would you drop these tables? Give **all** valid sequences.



Governor, Mayor, City, State

Mayor, Governor, City, State

Mayor, City, Governor, State

Mayor, Governor, State, City

... ..

Problem 4: Foreign keys

- Consider relation schemas below, with primary keys underlined.
 - City (name, state, population, elevation)
 - State (name, region)
 - Mayor (name, city, state, party)
 - Governor (name, state, party)
- Each **mayor governs** exactly one **city**. Each **governor governs** exactly one **state**.
- (a)** Write create table statements that encode these relation schemas and business rules with the right foreign key constraints.

```
create table City (  
  name      varchar(64),  
  state     varchar(32),  
  population number,  
  elevation  number,  
  primary key (name, state)  
  foreign key (state) references State(name) on delete cascade  
);
```


```
create table Mayor (  
  name varchar(128),  
  city  varchar(64),  
  state varchar(32),  
  party varchar(32),  
  primary key (name, city, state),  
  foreign key (state) references State(name) on delete cascade  
  foreign key (city) references City(name) on delete cascade  
);
```

```
create table State (  
  name varchar(32) primary key,  
  region varchar(32)  
);
```

```
create table Governor (  
  name varchar(128),  
  state varchar(32),  
  party varchar(32),  
  primary key (name, state),  
  foreign key (state) references State(name)  
  on delete cascade  
);
```


Problem 4: Foreign keys

- Consider relation schemas below, with primary keys underlined.
 - City (name, state, population, elevation)
 - State (name, region)
 - Mayor (name, city, state, party)
 - Governor (name, state, party)
- Each mayor governs exactly one city. Each governor governs exactly one state.
- (b)** In what order would you drop these tables? Give **all** valid sequences.



Governor, Mayor, City, State

Mayor, Governor, City, State

Mayor, City, Governor, State