

INFO 210: Database Management Systems

Topic 1

Introduction to databases
Overview of database design

supplementary material:
“A First Course in Database Systems” Ch. 1

Topic 1 outline

- Part 1: What is a DBMS?
- Part 2: Introducing data models
- Part 3: Introducing SQL
- Part 4: Sets
- Part 5: Summary

Part 1: What is a DBMS?

- A **database** is a large integrated collection of data
 - Models a real-world enterprise
 - Augments raw data with metadata, to give meaning to the data
- A **Database Management System (DBMS)** is a software package designed to store and manage databases
- Our focus: **Relational Database Management Systems (RDBMS)**, centered around the **relational model**
 - Entities (e.g., students and courses)
 - Relationships (e.g., students taking courses)

The role of a DBMS

- Serves as an intermediary between the user and the database
- Enables the sharing of data
- Supports multiple alternative views of the data

A traditional database application

- Build a system to store and access information about
 - students
 - courses
 - professors
 - who takes what, who teaches what
- Functionality
 - record enrollment information
 - compute GPA for each student after each term
 - analyze student enrollment and performance in different courses, majors, departments etc

Can we do this without a DBMS?

- Sure we can! We can store data in files:
 - *students.txt* *courses.txt* *professors.txt*



A programmer can write a C or Java program to implement specific tasks.

Different users can execute these tasks (run the programs)

Without a DBMS


- Task: Enroll “Mary” in “INFO210”
 - write a Java program that does the following:

```
read students.txt  
read courses.txt  
find & update record “Mary”  
find & update record “INFO210”  
write students.txt  
write courses.txt
```

Without a DBMS

- System crashes

read *students.txt*
read *courses.txt*
find & update record “Mary”
find & update record “INFO210”
write *students.txt*
write *courses.txt*



- Large data sets (say, 50GB)
- Simultaneous access by many users


Without a DBMS

- Format of *students.txt* changes
- Certain users should only be allowed to see parts of the file *courses.txt*

```
read students.txt  
read courses.txt  
find & update record "Mary"  
find & update record "INFO210"  
write students.txt  
write courses.txt
```

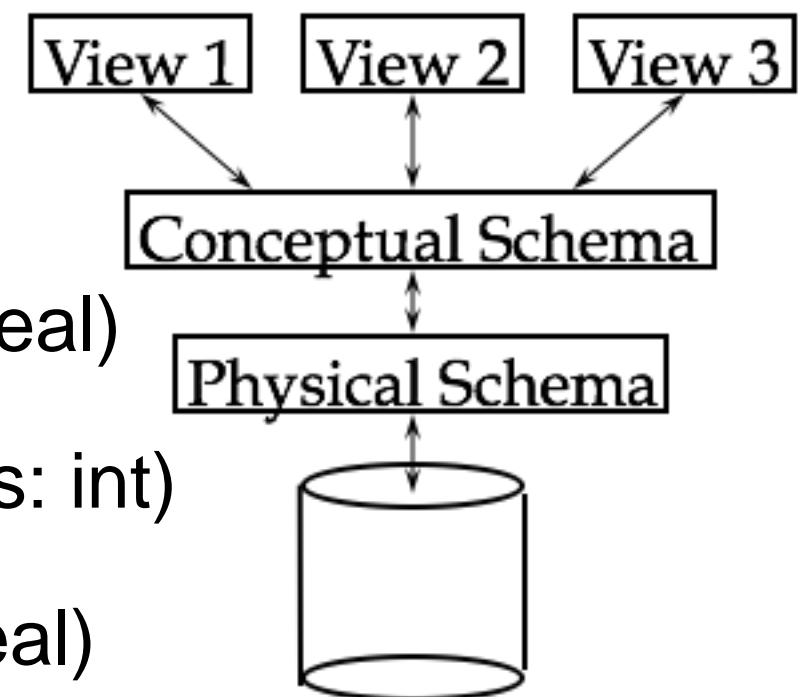
Enter a DBMS

- When in doubt - introduce a level (or levels) of abstraction!
- Many **views (external schemas)**, one **conceptual (logical) schema**, one **physical schema**
- Views describe how users see the data.
- A conceptual schema defines the logical structure.
- A physical schema describes the files and indexes used.



Example


- External schema (view)
 - Course_Info(cid: string, enrollment: int)
- Conceptual schema
 - Students (sid:string, name: string, gpa: real)
 - Courses (cid: string, name: string, credits: int)
 - Enrolled (sid: string, cid: string, grade: real)
- Physical schema
 - Courses, Enrolled stored as unordered files
 - Students stored sorted by sid



A key concept: data independence



- Applications (and users) are insulated from how data is structured and stored
- **Logical data independence**: protection from changes in logical structure of the data
- **Physical data independence**: protection from changes in physical structure of the data



So, why use a DBMS?

- Technical reasons
 - Data independence and efficient access
 - Reduced application development time
 - Data integrity and security
 - Uniform data administration
 - Concurrent access, recovery from crashes
- Business reasons
 - Reusing existing approaches makes data management more cost-effective!

A word of caution

- DBMS give us the tools to make data management more convenient, efficient, cost-effective
- But, like any technology, databases will only make our lives easier when used appropriately!

Part 2: Data models

- A **data model** is a collection of concepts for describing data.
- A **schema** is a description of a particular collection of data, using a given data model.
- We will look at two data models in this course
 - **relational model** - used to design logical and external schemas in relational databases
 - **entity-relationship (ER)** model - used for conceptual design

ER model basics

- Introduced by Chen in 1976.
- An **entity** is a real-world object distinguishable from other objects. Entities are described using a set of **attributes**.
- An **attribute** has a name and a datatype.
- An **entity set** is a collection of similar entities: all have the same set of attributes.
- A **relationship** is an association among 2 or more entities.
 - can be 1-to-1, 1-to-many, many-to-many
 - we can also express other interesting **constraints** on the relationships

The relational model

- Introduced by Edgar F. Codd in 1970 (Turing award)
- At the heart of relational database systems
 - the basic abstraction is a *relation* (a table)
 - *tuples* are stored in rows
 - *attributes* of tuples are stored in columns
 - conceptually, a relation is a *set* of tuples
- Why this model?
 - Simple yet powerful
 - Great for processing very large data sets in bulk

Some terminology

- The relational model is implemented (with some variations) by several RDBMS
 - IBM DB2, Oracle, Sybase, Microsoft SQL Server, MySQL, Postgress,
- Conceptually: relations are **sets of tuples**
- Reality: relations are implemented as **bags of tuples** (multisets).
- **Keys** allow us to unambiguously refer to tuples.

Relations

- Relations are used to describe two concepts
 - **Entities**, e.g., student, course, department
 - **Relationships**, e.g., student-enrolled-in-course, course-offered-by-department

Students

sid	name	GPA
1234	Joe	3.2
5678	Ann	4.0

key

Departments

did	name
1	INFO
2	MATH

Courses

cid	did	name	credits
110	1	HCI	3
210	1	Databases	3
312	2	Linear Algebra	3

Enrollment

did	cid	sid	term	grade
1	110	1234	SP11	A
2	312	1234	SP11	B
1	210	5678	FA12	A-
1	210	1234	FA12	B+

An alternative schema

Students

sid	name	GPA
1234	Joe	3.2
5678	Ann	4.0

Departments

did	name
1	INFO
2	MATH

Enrollment

did	sid	cid	name	credits	term	grade
1	1234	110	HCI	3	SP11	A
2	1234	312	Linear Algebra	3	SP11	B
1	5678	210	Databases	3	FA12	A-
1	1234	210	Databases	3	FA12	B+

Is this a well-designed
schema?
I have one word for you: *normalization!*

Part 3: SQL

- **SQL** (“seekwel”) stands for “Structured Query Language”
- Made up of 2 parts:
 - Data Definition Language (**DDL**) - used to create or modify a relational schema
 - Data Manipulation Language (**DML**) - used to retrieve or modify data in a schema
- We call SQL statements - **queries**

Creating relations (DDL)

```
create table XXX(  
  attribute dataType [keywords],  
  ...  
);
```

```
create table Students (  
  sid number primary key,  
  name varchar(128) not null,  
  gpa number  
);
```

```
create table Departments (  
  did number primary key,  
  name varchar(128) not null  
);
```

```
create table Courses (  
  cid number,  
  did number,  
  name varchar(128) unique,  
  credits number default 3,  
  primary key (cid, did),  
  foreign key (did) references Departments(did)  
);
```

Students

sid	name	GPA
1234	Joe	3.2
5678	Ann	4.0

key

Primary keys in a relation make sure that tuples can be distinguished from one another

Courses

cid	did	name	credits
110	1	HCI	3
210	1	Databases	3
110	2	Linear Algebra	3

Foreign keys enforce **referential integrity**

Creating relations (DDL)

```
create table Enrollment (  
  did number,  
  cid number,  
  sid number,  
  term varchar(32),  
  grade char(2),  
  primary key (cid, did, sid),  
  foreign key (cid, did) references Courses(cid, did),  
  foreign key (sid) references Students(sid)  
);
```

Students

sid	name	GPA
1234	Joe	3.2
5678	Ann	4.0

Courses

cid	did	name	credits
110	1	HCI	3
210	1	Databases	3
312	2	Linear Algebra	3


Enrollment

did	cid	sid	term	grade
1	110	1234	SP11	A
2	312	1234	SP11	B
1	210	5678	FA12	A-
1	210	1234	FA12	B+

Dropping relations (DDL)

```
drop table XXX;
```

```
drop table Enrollment;  
drop table Courses;  
drop table Departments;  
drop table Students;
```




why dropped in this particular order?

is this the only possible order?

Dropping relations (DDL)

```
drop table XXX;
```

```
drop table Enrollment;  
drop table Courses;  
drop table Departments;  
drop table Students;
```



why dropped in this particular order?

is this the only possible order?

Accessing data: queries (DML)

```
select XXX  
from XXX;
```

```
select *  
from Students;
```

```
select *  
from Students  
where gpa > 3.0;
```

```
select name  
from Students  
where gpa > 3.0;
```

```
update Students  
set name = 'Mike'  
where sid = 4;
```

```
select *  
from Enrollment  
where grade is null;
```

```
select did  
from Enrollment;
```

```
select distinct did  
from Enrollment;
```

Accessing data: queries (DML)

this is a *join*

```
select *  
from Students, Enrollment  
where Students.sid = Enrollment.sid  
and Enrollment.did = 100;
```

Accessing data: queries (DML)

these queries use *aggregation*

```
select count(*)  
from Students;
```

```
select min(gpa), max(gpa)  
from Students;
```

```
select term, count(*), avg(grade)  
from Enrollment  
group by term;
```

A key concept: SQL is declarative



- SQL is a declarative language: we say *what* we want to do, not *how* to do it
- This is important for two reasons:
 1. usability
 2. efficiency

Populating relations (DML)

```
insert into Students (sid, name, GPA) values (1, 'Jane', 4);
```

```
insert into Departments values (100, 'Math');
```

```
insert into Departments (name, did) values ('CIS', 200);
```

```
insert into Departments (name) values ('Italian');
```

do all of these work?

NB: SQL is not case-sensitive

NB: We use single quotes to quote strings

Part 4: Sets

- This course requires you to have a solid understanding of sets.
- Homework 1 has a sets section, it is your responsibility to refresh your memory on this material (or to learn it if you are not familiar with it)
- A quick refresher follows

Overview of sets

- A good overview at
- http://en.wikipedia.org/wiki/Set_theory#Basic_concepts
- A set is an *unordered collection* of objects
- An object belonging to a set is an *element* (or a *member*) of that set, written $a \in A$
- We denote sets with capital letters, elements with lowercase letters
- Examples of sets

Overview of sets (II)


- An empty set is a set that contains no elements. $A = \emptyset$
- We say that A is a *subset* of set B if all elements of A are also members of B .
Then B is a *superset* of A .
 $A \subseteq B$
 $B \supseteq A$
- A is a *proper subset* of B if A is a subset of B , and there exists at least one element $b \notin A$ such that $b \in B$.
 $A \subset B$
 $B \supset A$
- Any set A a subset of itself. A proper subset?
- $A = \emptyset$ is a subset of any set. A proper subset?

Overview of sets (III)

- The size of a set, denoted $|A|$, is the number of elements in the set
- The size of an empty set is 0
- Some sets are of infinite size, e.g., the set of all integers, all prime numbers, etc


Set operations

Intersection




$$A \cap B$$

Union



$$A \cup B$$

Difference




$$B \setminus A$$

Venn diagrams from

http://en.wikipedia.org/wiki/Venn_diagram

Set operations (II)

How can you express *symmetric difference* using other set operations?



Cartesian product

- Cartesian product of A and B, $A \times B$ denoted is a set of ordered pairs (a, b), where $a \in A, b \in B$

Example:

$$A = \{1, 2, 3\} \quad B = \{3, 4\}$$

$$A \times B = \{(1, 3), (1, 4), (2, 3), (2, 4), (3, 3), (3, 4)\}$$

Part 5: Summary

- Take-home message: Databases are cool
- Differences between databases and file systems
- The relational model
- The ER model
- Example of a database schema
- Examples of SQL queries
- A key concept: data independence
- A key concept: SQL is declarative

Useful abbreviations

- DB - database
- DBMS - database management system
- RDBMS - relational database management system
- DBA - database administrator
- SQL - structured query language
 - DDL - data definition language
 - DML - data manipulation language
- ER - entity-relationship
- ERM - entity-relationship model